

Bingo Spatial Data Prefetcher in Scarab

Abstract

This project implements Bingo, the HPCA 2019 spatial data prefetcher, in Scarab. Bingo learns spatial footprints within 2 KiB memory regions and uses a single pattern-history table (PHT) to balance coverage and accuracy. The implementation is adapted from the public ChampSim reference code while using Scarab's native prefetch queue, parameter, and statistics interfaces. On 10 course SPEC traces, Bingo improves geometric-mean IPC by 4.37% over a no-data-prefetch baseline and reduces unified L1 demand misses by 63.73%.

Technique

Bingo records the first access to a region in a 64-entry filter table. A second distinct block access promotes the region to a 128-entry accumulation table, where the footprint is collected. When the region generation ends, Bingo trains a 16K-entry, 16-way PHT. The lookup first attempts a precise PC+Address match. If it misses, entries with the same PC+Offset event vote on the predicted footprint using the upstream 20% threshold. This fallback improves coverage without maintaining two independent pattern tables.


Scarab Integration

The public Bingo code is a ChampSim LLC prefetcher. Scarab exposes a unified L1 path above its private dcache, so this implementation attaches Bingo to that unified L1 path. Demand accesses call the Bingo access hook and predictions are submitted through `pref_addto_ul1req_queue_set()`. A new unified-cache eviction callback ends region generations and trains the PHT. Configuration parameters and counters are registered with Scarab's existing generated parameter and statistics systems.

The implementation keeps the reference defaults: 2 KiB regions, 64 filter entries, 128 accumulation entries, 16K PHT entries, 16 PHT ways, 16-bit PC signatures, and a 20% voting threshold. The baseline disables all data prefetchers. The Bingo configuration changes only `--pref_bingo_on`.

Methodology

The evaluation uses the 10 SPEC traces provided by the course Docker image. Each workload is simulated twice, once with no data prefetcher and once with Bingo. Each run uses a 20M-instruction trace window with a 10M-instruction full-warmup interval. The scripts preserve raw Scarab output, parse the regular post-warmup statistics CSV files, and generate the IPC-speedup plot below.



Results

The geometric-mean IPC speedup is 1.0437. The strongest improvement is 523.xalancbmk_r at 1.2849, followed by 500.perlbench_r at 1.0916 and 520.omnetpp_r at 1.0604. Two workloads regress slightly: 502.gcc_r reaches 0.9975 and 557.xz_r reaches 0.9868.

Across all workloads, unified L1 demand misses fall from 633,345 to 229,697. Bingo queues 1,067,480 prefetch requests, with no queue rejection in this experiment. Scarab records 418,703 useful prefetch hits. The observed hit-to-queued ratio is 39.22%; it is an implementation-level usefulness ratio, not a direct substitute for the paper's accuracy metric.

Limitations

This is a Scarab mechanism reproduction, not a bit-identical replay of the paper's ChampSim experiments. The cache attachment point, simulator timing model, trace set, and simulation length differ from the paper. The current Windows Scarab checkout also expects newer architecture-register trace metadata, while the course traces use the older Lab 2 format. The formal evaluation therefore runs the same Bingo source changes on the Lab-2-compatible course Scarab base inside Docker.

Reproduction

Run scripts/run_bingo_spec.py for the two-configuration SPEC experiment and scripts/plot_bingo_results.py for the metrics CSV and figure. Full commands, paths, build instructions, and generated artifacts are documented in README_BINGO.md.